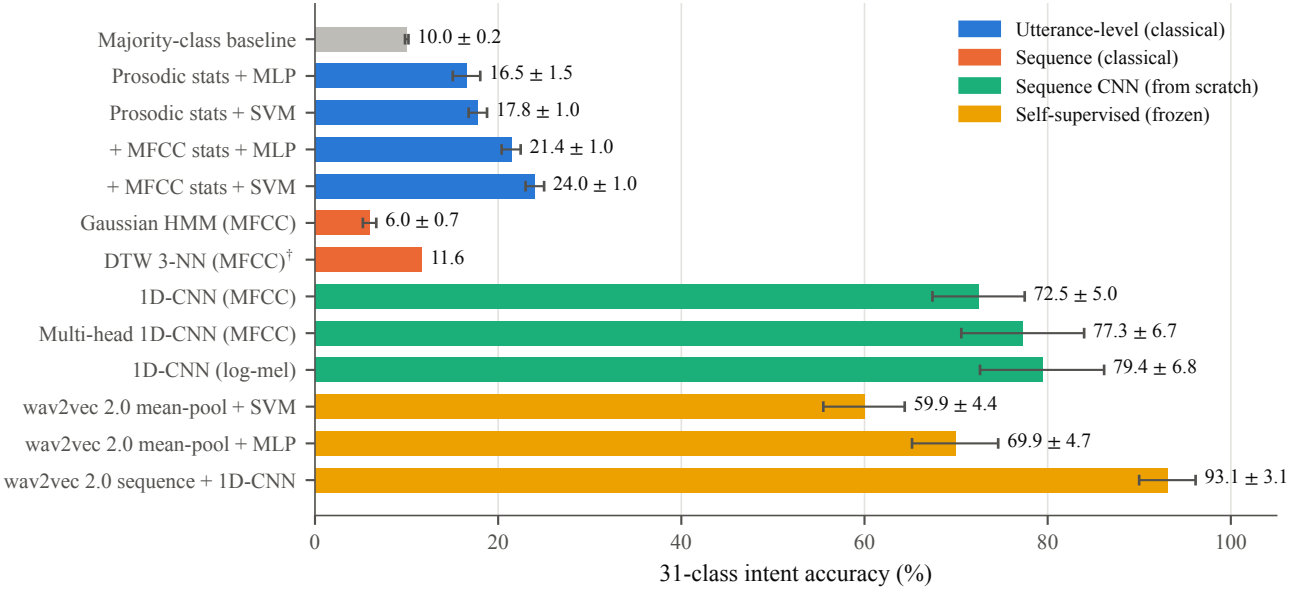
Utterance-level (classical)
Sequence (classical)
Sequence CNN (from scratch)
Self-supervised (frozen)
Majority-class baseline
10.0 ± 0.2
Prosodic stats + MLP
16.5 ± 1.5
Prosodic stats + SVM
17.8 ± 1.0
+ MFCC stats + MLP
21.4 ± 1.0
+ MFCC stats + SVM
24.0 ± 1.0
Gaussian HMM (MFCC)
6.0 ± 0.7
DTW 3-NN (MFCC)†
11.6
1D-CNN (MFCC)
72.5 ± 5.0
Multi-head 1D-CNN (MFCC)
77.3 ± 6.7
1D-CNN (log-mel)
79.4 ± 6.8
wav2vec 2.0 mean-pool + SVM
59.9 ± 4.4
wav2vec 2.0 mean-pool + MLP
69.9 ± 4.7
wav2vec 2.0 sequence + 1D-CNN
93.1 ± 3.1
0
20
40
60
80
100
31-class intent accuracy (%)