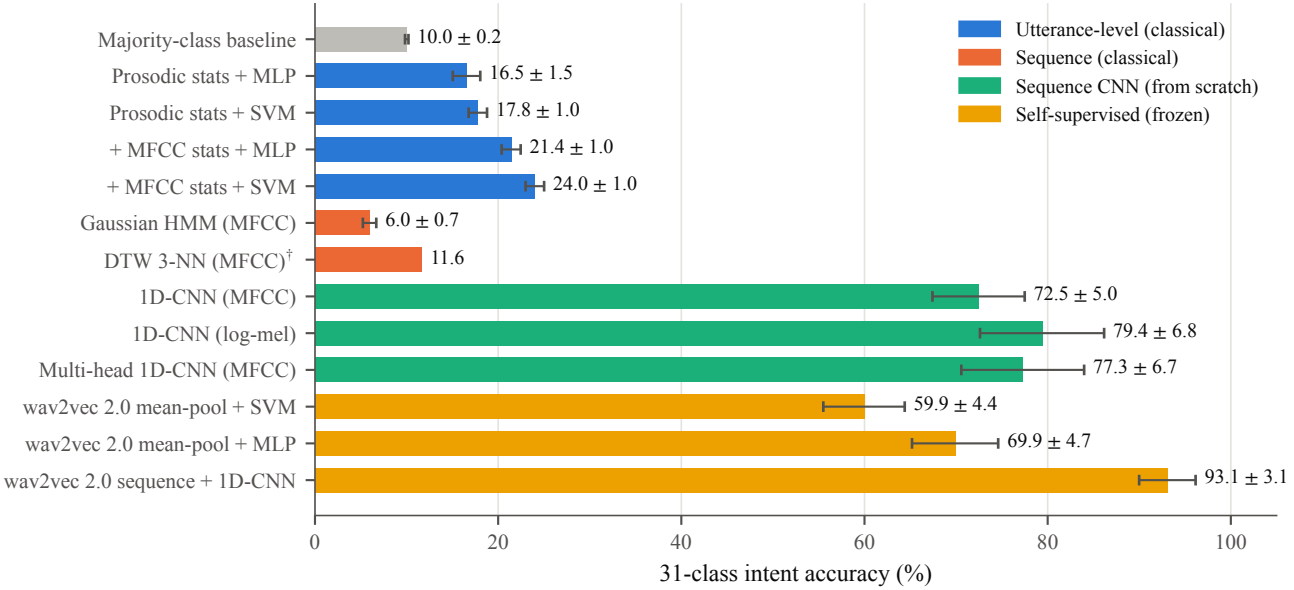
Utterance-level (classical)
Sequence (classical)
Sequence CNN (from scratch)
Self-supervised (frozen)
Majority-class baseline
Prosodic stats + MLP
Prosodic stats + SVM
+ MFCC stats + MLP
+ MFCC stats + SVM
Gaussian HMM (MFCC)
DTW 3-NN (MFCC)†
1D-CNN (MFCC)
1D-CNN (log-mel)
Multi-head 1D-CNN (MFCC)
wav2vec 2.0 mean-pool + SVM
wav2vec 2.0 mean-pool + MLP
wav2vec 2.0 sequence + 1D-CNN
10.0 ± 0.2
16.5 ± 1.5
17.8 ± 1.0
21.4 ± 1.0
24.0 ± 1.0
6.0 ± 0.7
11.6
72.5 ± 5.0
79.4 ± 6.8
77.3 ± 6.7
59.9 ± 4.4
69.9 ± 4.7
93.1 ± 3.1
0
20
40
60
80
100
31-class intent accuracy (%)